
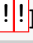


Deep_Learning_Assignment

November 13, 2023

1 Sentiment Analysis With Transformers

Name: Tarun Kumar Reddy

```
[23]:  pip install -q rouge-score  
 pip install -q git+https://github.com/keras-team/keras-nlp.git --upgrade
```

```
[23]: [' Installing build dependencies ... \x1b[?25l\x1b[?25hdone',  
      ' Getting requirements to build wheel ... \x1b[?25l\x1b[?25hdone',  
      ' Preparing metadata (pyproject.toml) ... \x1b[?25l\x1b[?25hdone',  
      '\x1b[?25l      \x1b[90m                        \x1b[0m  
\x1b[32m0.0/950.8 kB\x1b[0m \x1b[31m?\x1b[0m eta \x1b[36m-:--:--\x1b[0m',  
      '\x1b[2K      \x1b[91m      \x1b[0m\x1b[91m \x1b[0m\x1b[90m  
      \x1b[0m \x1b[32m307.2/950.8 kB\x1b[0m \x1b[31m9.1 MB/s\x1b[0m eta  
\x1b[36m0:00:01\x1b[0m',  
      '\x1b[2K      \x1b[90m                        \x1b[0m  
\x1b[32m950.8/950.8 kB\x1b[0m \x1b[31m17.2 MB/s\x1b[0m eta  
\x1b[36m0:00:00\x1b[0m',  
      '\x1b[?25h\x1b[?25l      \x1b[90m                        \x1b[0m  
\x1b[32m0.0/6.5 MB\x1b[0m \x1b[31m?\x1b[0m eta \x1b[36m-:--:--\x1b[0m',  
      '\x1b[2K  
\x1b[91m                        \x1b[0m\x1b[91m \x1b[0m  
\x1b[32m6.5/6.5 MB\x1b[0m \x1b[31m207.9 MB/s\x1b[0m eta \x1b[36m0:00:01\x1b[0m',  
      '\x1b[2K      \x1b[90m                        \x1b[0m  
\x1b[32m6.5/6.5 MB\x1b[0m \x1b[31m107.1 MB/s\x1b[0m eta \x1b[36m0:00:00\x1b[0m',  
      '\x1b[?25h Building wheel for keras-nlp (pyproject.toml) ...  
\x1b[?25l\x1b[?25hdone']
```

1.1 Import Required Libraries

```
[26]: import tensorflow as tf  
      from tensorflow import keras  
      from tensorflow.keras import layers  
      import keras_nlp
```

Using TensorFlow backend

1.2 Loading Data

```
[27]: vocab_size = 20000
maxlen = 150

(x_train, y_train), (x_val, y_val) = keras.datasets.imdb.
    ↳load_data(num_words=vocab_size)

print(len(x_train), "Training sequences")
print(len(x_val), "Validation sequences")

x_train = keras.utils.pad_sequences(x_train, maxlen=maxlen)
x_val = keras.utils.pad_sequences(x_val, maxlen=maxlen)
```

25000 Training sequences
25000 Validation sequences

1.3 Building Architecture

```
[29]: embed_dim = 256
num_heads = 8
ff_dim = 32
intermediate_dim = 512
num_layers = 3

[31]: input_ = keras.Input(shape=(None,), dtype="int32")

x = keras_nlp.layers.TokenAndPositionEmbedding(vocabulary_size=vocab_size,
                                                sequence_length=maxlen,
                                                embedding_dim = embed_dim,
                                                mask_zero=True,)(input_)

for _ in range(num_layers):
    encoder_layer = keras_nlp.layers.TransformerEncoder(num_heads= num_heads,
                                                         ↳
    ↳intermediate_dim=intermediate_dim,)
    x = encoder_layer(x)

x = layers.GlobalAveragePooling1D()(x)
x = layers.Dropout(rate=0.2)(x)
x = layers.Dense(30, activation="relu")(x)
x = layers.Dense(20, activation="relu")(x)

output_ = layers.Dense(2, activation="softmax")(x)

model = keras.Model(input_, output_)
```

```
[32]: model.summary()
```

Model: "model_2"

Layer (type)	Output Shape	Param #
input_4 (InputLayer)	[(None, None)]	0
token_and_position_embedding_2 (TokenAndPositionEmbedding)	(None, None, 256)	5158400
transformer_encoder (TransformerEncoder)	(None, None, 256)	527104
transformer_encoder_1 (TransformerEncoder)	(None, None, 256)	527104
transformer_encoder_2 (TransformerEncoder)	(None, None, 256)	527104
global_average_pooling1d_2 (GlobalAveragePooling1D)	(None, 256)	0
dropout_14 (Dropout)	(None, 256)	0
dense_18 (Dense)	(None, 30)	7710
dense_19 (Dense)	(None, 20)	620
dense_20 (Dense)	(None, 2)	42

=====
Total params: 6748084 (25.74 MB)
Trainable params: 6748084 (25.74 MB)
Non-trainable params: 0 (0.00 Byte)
=====

1.4 Training the model

```
[33]: model.compile(optimizer="adam", loss="sparse_categorical_crossentropy",  
    ↪ metrics=["accuracy"])  
history = model.fit(x_train, y_train, batch_size=32, epochs=5,  
    ↪ validation_data=(x_val, y_val))
```

Epoch 1/5

782/782 [=====] - 88s 100ms/step - loss: 0.6960 -


```
accuracy: 0.5049 - val_loss: 0.6934 - val_accuracy: 0.5000
Epoch 2/5
782/782 [=====] - 57s 73ms/step - loss: 0.6934 -
accuracy: 0.4992 - val_loss: 0.6932 - val_accuracy: 0.5000
Epoch 3/5
782/782 [=====] - 55s 71ms/step - loss: 0.6934 -
accuracy: 0.5007 - val_loss: 0.6932 - val_accuracy: 0.5000
Epoch 4/5
782/782 [=====] - 46s 59ms/step - loss: 0.6932 -
accuracy: 0.4999 - val_loss: 0.6933 - val_accuracy: 0.5000
Epoch 5/5
782/782 [=====] - 46s 59ms/step - loss: 0.6932 -
accuracy: 0.5003 - val_loss: 0.6932 - val_accuracy: 0.5000
```

```
[34]: keys = ['accuracy', 'val_accuracy']
progress = {k:v for k,v in history.history.items() if k in keys}

import pandas as pd
pd.DataFrame(progress).plot()

plt.xlabel("epochs")
plt.ylabel("accuracy")

plt.grid(True)
plt.show()
```



1.5 Evaluating the model using “Precision”, “Recall” and “F1-score”

```
[35]: import numpy as np
from sklearn import metrics
from sklearn.metrics import confusion_matrix
from sklearn.metrics import ConfusionMatrixDisplay
import matplotlib.pyplot as plt
y_pred = np.argmax(model.predict(x_val), axis=1)

print("Classification report for classifier %s:\n%s\n"
      % (model, metrics.classification_report(y_val, y_pred)))

confMatrix = confusion_matrix(y_true = y_val, y_pred = y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix = confMatrix)
disp.plot()
plt.show()
```

782/782 [=====] - 16s 18ms/step

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:

UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

Classification report for classifier <keras.src.engine.functional.Functional object at 0x7d1afd441720>:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	12500
1	0.50	1.00	0.67	12500
accuracy			0.50	25000
macro avg	0.25	0.50	0.33	25000
weighted avg	0.25	0.50	0.33	25000

